Financial Literacy and Decision Support: Using LLMs to Drive SDG Progress

Abstract

This project focuses on the development of an industry-specific conversational Large Language Model (LLM) tailored for the finance sector. Leveraging pre-trained models from Hugging Face, the project fine-tunes the model using domain-specific datasets. Advanced fine-tuning techniques, including Parameter-Efficient Fine-Tuning (PEFT), Low-Rank Adaptation (LoRA), and Quantized LoRA (QLoRA), are utilized to optimize training efficiency while maintaining high performance. Implementation and interaction with the fine-tuned model are carried out using Google Colab, which provides an accessible platform for development and demonstration. This project underscores the importance of domain-specific AI solutions in addressing complex challenges in the finance industry, offering significant potential to enhance user experiences and improve decision-making processes. Future expansions, including broader datasets and real-time financial data integration, could further increase the model's capabilities and relevance.

1. Introduction

The rapid advancements in Artificial Intelligence (AI) have unlocked immense potential across various industries. Among these, the finance sector has emerged as a critical area for leveraging AI technologies due to its reliance on complex data analysis, decision-making, and customer interaction. Large Language Models (LLMs) have demonstrated exceptional capabilities in understanding and generating human-like language, making them suitable for addressing domain-specific challenges.

This project aims to develop a conversational LLM tailored specifically for the finance industry. By fine-tuning a pre-trained LLM from Hugging Face on finance-focused datasets, the model is equipped to provide accurate, contextually relevant responses to industry-specific queries. The choice of the finance sector is driven by its dynamic nature, the critical need for precision, and the vast availability of structured and unstructured data.

The project employs advanced fine-tuning techniques such as PEFT, LoRA, and QLoRA to ensure computational efficiency while achieving high model performance. Google Colab serves as the primary environment for implementing and testing the conversational capabilities of the model, allowing users to engage with it directly. This report outlines the end-to-end process, from data collection and preprocessing to model fine-tuning and evaluation, showcasing the development of an intelligent, domain-specific conversational agent.

The outcomes of this project highlight the potential of LLMs in transforming industry-specific communication, offering insights into how they can enhance decision-making, customer interaction, and overall operational efficiency in the finance sector.

1.1 Background

The United Nations 2030 Agenda for Sustainable Development aims to address global challenges like poverty, hunger, inequality, and climate change. However, the latest report highlights significant progress hurdles, with many targets falling behind schedule.

Technology, particularly AI and LLMs, can play a crucial role in accelerating progress towards these goals. By improving financial inclusion, these technologies can empower individuals and communities, leading to economic growth and poverty reduction. AI can enhance financial literacy, support data-driven decision-making, and optimize resource allocation. Additionally, it can revolutionize industries like finance, education, and healthcare.

To maximize the positive impact of AI and LLMs, it's crucial to prioritize their ethical and responsible use. By addressing potential biases and ensuring equitable access, we can harness the power of technology to drive sustainable development and create a more equitable future.

1.2 Problem Statement

The Digital Divide in Financial Literacy

A significant portion of the global population remains financially illiterate, lacking a basic understanding of financial concepts. This lack of knowledge hinders financial empowerment, particularly among vulnerable groups like women and low-income individuals. While governments are striving to expand access to financial services, it's imperative to address the underlying issue of financial literacy.

The Role of Technology in Bridging the Gap

Financial technology offers a promising solution to this challenge. By leveraging digital platforms and mobile apps, financial institutions can reach underserved populations and provide accessible financial services. These platforms can offer personalized financial advice, interactive learning tools, and simplified financial products, making it easier for individuals to manage their finances effectively.

Empowering Individuals Through Financial Literacy

Through financial education, individuals can make informed decisions about savings, investments, and debt management. This empowers them to build financial resilience, improve their standard of living, and contribute to economic growth. As financial technology continues to evolve, it has the potential to revolutionize financial literacy and inclusion, creating a more equitable and prosperous future for all

2. Project Goals

Our goal is to develop an AI-powered financial advisor specifically designed to empower underserved communities. By leveraging the capabilities of large language models, we aim to bridge the financial literacy gap and provide accessible, personalized financial guidance.

The Power of AI

AI's ability to process vast amounts of data quickly and accurately can revolutionize financial services. By analyzing individual financial situations, AI can provide tailored recommendations for:

- Investment Strategies**:** Identifying suitable investment options based on risk tolerance and financial goals.
- Debt Management**:** Optimizing debt repayment strategies to minimize interest costs.
- Budgeting and Savings**:** Analyzing spending patterns and suggesting savings strategies.
- Fraud Detection**:** Proactively monitoring financial transactions for suspicious activity.

Benefits for Underserved Communities

- **Accessibility:** AI-powered tools can be accessed anytime, anywhere, making financial advice more accessible to underserved populations.
- **Affordability:** AI-driven solutions can be more cost-effective than traditional financial advisors, making them accessible to a wider range of individuals.
- **Personalized Guidance:** AI can provide tailored financial advice based on individual needs and circumstances.
- **Financial Literacy:** AI can help individuals understand complex financial concepts and make informed decisions.

Building on Existing Models

Models like BloombergGPT and FinGPT demonstrate the potential of AI in the financial sector. By leveraging these advancements, we can create a financial LLM that is specifically designed to address the needs of underserved communities.

By developing this AI-powered financial advisor, we can help individuals achieve financial stability and security, ultimately contributing to a more equitable and prosperous society

2.2 Alignment with SDGs

This project strongly aligns with two Sustainable Development Goals (SDGs):
SDG 1 (No Poverty) and SDG 8 (Decent Work and Economic Growth).

Under SDG 1, it directly contributes to Target 1.4, which aims to ensure that all individuals, particularly the poor and vulnerable, have access to financial resources and services, including microfinance. By using a financial Large Language Model (LLM) tailored to underserved communities, the project empowers individuals with personalized financial tools and education to improve their economic situation. Enhanced access to financial literacy reduces the likelihood of debt traps, fosters better savings habits, and encourages financial independence. For those who are excluded from traditional banking systems, this project acts as a gateway to financial empowerment, enabling them to build wealth and actively work towards reducing poverty.

Similarly, the project supports SDG 8, with a focus on Target 8.10, which emphasizes expanding access to financial services and strengthening the capacity of domestic financial institutions. By offering AI-powered financial guidance, the project enables individuals and small businesses to make informed decisions regarding investment, debt management, and savings, promoting greater participation in the formal economy. This fosters entrepreneurship, job creation, and economic growth while improving access to financial products such as loans and insurance. By enhancing financial literacy and inclusion, the project equips underserved populations to better navigate economic challenges and seize opportunities, ultimately contributing to inclusive economic growth and decent work for all. Through its focus on accessibility, empowerment, and inclusivity, this project aligns with global efforts to reduce poverty and promote sustainable economic development.

3. Methodology

3.1 Dataset

The fingpt-fiqa_qa dataset, part of the FinGPT initiative on Hugging Face, is a curated collection of approximately 17,100 financial question-and-answer pairs. Tailored for instruction-tuned large language models, this dataset addresses the specific complexities of the financial sector, including topics like market analysis, investment strategies, financial products, and regulatory compliance. With a focus on domain-specific terminology, numerical reasoning, and dynamic financial scenarios, it is designed to support applications requiring precision and contextual understanding, such as personalized financial advisory, sentiment analysis, and insight extraction from unstructured financial documents.

This dataset’s high-quality annotations and real-world relevance make it a powerful resource for building advanced AI tools for the financial industry. In this project, the dataset played a pivotal role in training the model to handle finance-specific queries effectively. Its diverse content, including a wide range of questions and instructional data, enhanced the model’s ability to provide accurate, context-aware responses to complex financial queries. By equipping the model with a strong grasp of financial jargon and concepts, the fingpt-fiqa_qa dataset significantly improved its performance, demonstrating its value as a benchmark for domain-specific NLP advancements.

3.2 Model selection

The Llama-2-7b-chat-hf model was chosen for this project due to its ability to handle complex financial queries with accuracy and contextual relevance. With 7 billion parameters, the model strikes a balance between computational efficiency and performance, making it ideal for resource-constrained setups. As a fine-tuned version of the Llama-2 base model, it excels in natural language understanding and maintaining context in extended interactions, critical for addressing the diverse and intricate nature of financial inquiries.

The model was fine-tuned using the FinGPT-fiqa_qa dataset, a specialized collection of financial questions and answers covering real-world scenarios, terminology, and problem-solving contexts. This process enhanced the model's understanding of financial jargon, market trends, and complex queries. Llama-2's open-source availability through Hugging Face ensures flexibility for future enhancements, making it a robust choice for domain-specific applications.

4. System Architecture and Training Details

System Architecture Overview

The system architecture for fine-tuning a financial language model is designed to streamline data preparation and model training. Key components include:

1. **Dataset Preparation**:
    - The **FinGPT-fiqa_qa dataset**, comprising financial question-answer pairs, is used for fine-tuning.
    - Preprocessing involves merging columns (instruction, input, and output) into a unified text format separated by a unique delimiter ###. This format emulates a conversational prompt-response structure.
    - After preprocessing, unnecessary columns are removed, and the data is split into training (80%) and testing (20%) sets. The training set is shuffled to ensure unbiased learning.
2. **Model Selection and Setup**:
    - The **Llama-2-7b-chat-hf** model is chosen for its ability to manage conversational tasks efficiently.
    - Fine-tuning creates a new model named "finance-chatbot," tailored to understand financial queries and provide accurate, context-aware responses.

3. **Training Environment and Configuration**:
    - **QLoRA (Quantized Low-Rank Adaptation)** is used to enhance memory efficiency, with 4-bit quantization (bnb_4bit and nf4) reducing resource usage.
    - Hugging Face's **Training Arguments** manages settings like batch size, learning rate, and epochs. A cosine learning rate scheduler ensures smooth convergence.
4. **Model Fine-Tuning**:
    - Fine-tuning leverages **Parameter-Efficient Fine-Tuning (PEFT)** using **LoRA** to update only a subset of parameters, reducing computational load.
    - Text is tokenized into a suitable format, with special tags (<s>[INST] and [/INST]) marking prompts and responses.
    - The **SFTTrainer** class executes supervised fine-tuning, training the model to align with financial domain queries.
5. **Post-Training Evaluation**:
    - After fine-tuning, the model is evaluated on unseen data to verify its accuracy and relevance. Metrics like response quality and ability to handle new financial queries are assessed.
    - The final model, "finance-chatbot," is saved for further deployment and enhancements.

Training Details

Model and Framework:

- **Llama-2-7b-chat-hf**, optimized for conversations, is fine-tuned to handle financial domain queries.
- The **SFTTrainer** API ensures efficient training, leveraging QLoRA for memory optimization on resource-constrained hardware.

Hyperparameters:

- Key configurations include a cosine learning rate scheduler, batch size, gradient accumulation, and LoRA-specific parameters like lora_r, lora_alpha, and lora_dropout.

Data Transformation:

- Financial question-answer pairs are structured into prompt-response text pairs.
- Special tokens (<s>[INST] and [/INST]) guide the model to distinguish between user inputs and generated outputs.

Training and Evaluation:

- The training process iteratively adjusts weights based on the transformed dataset, focusing on real-world financial scenarios.
- Post-training, the model is evaluated against test data to ensure generalization and accurate performance.

The system architecture and training details ensure that the financial language model is effectively fine-tuned to provide accurate, contextually relevant responses for financial queries. By leveraging advanced training techniques such as QLoRA and PEFT, the project makes it possible to train large models efficiently, enabling deployment for real-world applications in enhancing financial literacy and inclusion.

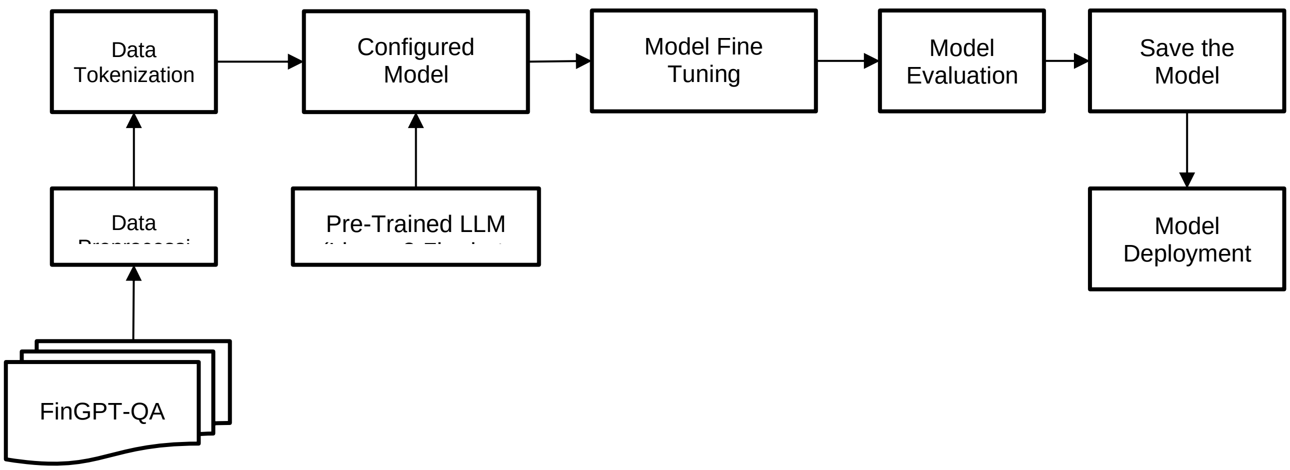

4. Analysis of Community Impact

4.1 Scalability

It is a critical aspect of ensuring that this solution can reach diverse communities with varying needs and resources. Key factors include:

- **Customizability for Different Communities:** The solution can be adapted to cater to varying levels of financial literacy, from basic concepts for beginners to advanced financial guidance for experienced users. This ensures inclusivity across a wide range of demographics.
- **Language and Regional Adaptation:** By fine-tuning the LLM with multilingual support and incorporating regional financial data, the system can address the unique financial challenges of different communities.

- **Technology Integration:** The use of lightweight models and cloud deployment enables the solution to scale efficiently, serving urban populations with high-speed internet access and rural communities with limited digital infrastructure.

4.2 Affordability

It is essential for ensuring equitable access to financial literacy tools, especially for underserved communities. The solution achieves this by:

- **Open-Source Foundations:** Leveraging open-source LLMs like Llama-2 reduces initial development costs and provides a foundation for continuous community-driven improvements.
- **Cloud-Based Deployments:** Deploying the solution on cost-effective cloud platforms minimizes infrastructure costs for users and developers. For offline use cases, lightweight deployment on mobile devices using optimized models (e.g., with QLoRA) ensures accessibility without high-end hardware.
- **Freemium Models:** Offering a basic version for free, supplemented with premium features for advanced users or financial institutions, ensures widespread adoption while generating revenue for sustainability.

4.3 Sustainability

The long-term sustainability of this project relies on its ability to foster lasting improvements in financial literacy and economic participation. This is achieved through:

- **Behavioral Change:** By providing users with accessible and actionable financial advice, the solution encourages better savings habits, investment strategies, and informed decision-making, contributing to improved financial resilience over time.
- **Economic Empowerment:** Increased financial literacy enables individuals and small businesses to access credit, manage debt effectively, and participate in economic activities, driving community-level economic growth.
- **Collaboration with Stakeholders:** Partnerships with financial institutions, NGOs, and governments ensure the tool remains relevant, up-to-date, and well-integrated into existing financial inclusion initiatives.
- **Data-Driven Feedback:** Using analytics to monitor user behavior and outcomes provides insights for continuous improvement, ensuring the solution adapts to changing economic and social conditions.

5. Ethical Concerns and Mitigation Strategies

Ethical concerns are critical in the development and deployment of AI-powered solutions, especially in the finance sector, where sensitive data and decision-making are involved. Below are the key concerns for the financial literacy project using LLMs and the strategies to mitigate them effectively:

5.1 Bias in Financial Data

- **Concern**: Financial data used for training may reflect inherent biases, such as exclusion of regional financial contexts, underrepresentation of underserved communities, or systemic inequalities.
- **Impact**: Biased output could perpetuate financial disparities or provide inaccurate advice to certain demographics.
- **Mitigation Strategies**:
    - Incorporate diverse datasets covering a wide range of regional, cultural, and economic contexts.
    - Regularly test and audit the model for biases using real-world scenarios and edge cases.
    - Collaborate with financial experts from varied backgrounds to identify and address biases.

5.2 Privacy and Security Risks

- **Concern**: User interactions with the financial literacy tool may involve sensitive financial information, raising privacy and data security risks.
- **Impact**: Breaches or mishandling of data could lead to loss of trust, financial harm, or legal repercussions.
- **Mitigation Strategies**:
    - Implement end-to-end encryption to secure data during transmission.
    - Use anonymization techniques to ensure no personally identifiable information (PII) is stored.
    - Adhere strictly to global data protection regulations, such as GDPR and CCPA, to ensure compliance and user safety.

5.3 Transparency

- **Concern**: Users may not fully understand the limitations of the AI tool, leading to overreliance on its recommendations.
- **Impact**: Misinterpretation of the tool's advice could result in poor financial decisions or unrealistic expectations.
- **Mitigation Strategies**:

- Provide clear disclaimers about the AI's role as a supplementary advisor, not a replacement for professional financial advice.
- Offer detailed explanations of the AI's decision-making process to enhance user understanding and trust.
- Maintain transparency about the datasets and training processes used to develop the tool.

5.4 Accessibility

- **Concern**: Without careful design, the tool could unintentionally exclude certain groups, such as those with limited internet access or low digital literacy.
- **Impact**: This would undermine the goal of inclusivity and fail to serve the most underserved populations.
- **Mitigation Strategies**:
  - Optimize the tool for low-bandwidth environments and mobile platforms.
  - Design an intuitive user interface with multilingual support and user-friendly explanations.

6. Conclusion and Future Work

6.1 Summary of Achievements

This project demonstrated the potential of Large Language Models (LLMs) to advance financial literacy and inclusion, particularly for underserved communities. Key accomplishments include:

- Development of a domain-specific conversational financial model fine-tuned using the FinGPT-fiqa_qa dataset.
- Implementation of efficient training techniques such as PEFT, LoRA, and QLoRA, optimizing resource usage while maintaining high performance.
- Deployment of an accessible, AI-powered financial advisor capable of providing tailored guidance on budgeting, investments, debt management, and more.
- Alignment with Sustainable Development Goals (SDGs) by addressing financial literacy gaps and empowering individuals with knowledge and tools to achieve financial stability.

6.2 Next Steps

Building on the current progress, the following future directions will further enhance the project's impact and scalability:

1. **Dataset Expansion**:
   - Integrating more diverse and comprehensive datasets, including regional financial data, to improve the model’s contextual relevance and reduce biases.
2. **Multilingual Support**:

- Extending the tool to support multiple languages, enabling broader accessibility for non-English-speaking communities worldwide.

3. **Real-Time Data Integration**:
    - Incorporating real-time financial data and news to offer up-to-date guidance and market insights.
4. **Collaboration with NGOs and Financial Institutions**:
    - Partnering with organizations focused on financial inclusion to deploy the tool in communities with limited access to traditional financial services.
5. **Interactive Features**:
    - Adding functionalities like scenario simulation and interactive tutorials to make financial concepts easier to grasp.

References:

[1] United Nations, "Sustainable Development Goals," *United Nations*, [Online]. Available: https://sdgs.un.org/goals.

[2] Q. Wu, Y. Xu, T. Xiao, Y. Xiao, Y. Li, T. Wang, Y. Zhang, S. Zhong, Y. Zhang, W. Lu, and Y. Yang, "Surveying attitudinal alignment between large language models vs. humans towards 17 sustainable development goals," *arXiv preprint arXiv:2404.13885*, Apr. 2024. [Online]. Available: https://arxiv.org/pdf/2404.13885.

[3] Silveira, C., & Reis, L. (2022). *Sustainability in Software Engineering: A Design Science Research Approach*. Presented at the 8th International Scientific Conference ERAZ - Knowledge Based Sustainable Development, Prague, Czech Republic, July 1–3, 2022. DOI: 10.31410/ERAZ.2022.317

[4] Global Financial Literacy Excellence Center (GFLEC), "S&P Global FinLit Survey," [Online]. Available: https://gflec.org/initiatives/sp-global-finlit-survey/.

[5] H. Yang, X.-Y. Liu, and C. D. Wang, "FinGPT: Open-Source Financial Large Language Models," Columbia University and New York University (Shanghai). [Online]. Available: https://doi.org/10.48550/arXiv.2306.06031

[6] Y. Li, S. Wang, H. Ding, and H. Chen, "Large Language Models in Finance: A Survey," *arXiv preprint*, 2024. Available: https://doi.org/10.48550/arXiv.2311.10723